Oversampling with sample function

Asked 1 year, 4 months ago Active 1 year, 4 months ago Viewed 196 times

1

I would like to create a `mtcars` dataset where all cylinders have 100 observations. For that, I would sample with replacement the existing observations.

I have tried the following code that, for some reason, does not produce 300 observations.

```
library(data.table)
mtcars <- data.table(mtcars)
resampling <- list()
set.seed(3)

cyl <- sort(unique(as.character(mtcars$cyl)))
for (i in 1:length(cyl)){

  min_obs_cyl <- 100
  dat_cyl <- mtcars[cyl == as.numeric(cyl[i]) ]
  resampling[[  cyl[i]  ]] <- dat_cyl[sample(1:nrow(dat_cyl),
                                             size = (min_obs_cyl - nrow(mtcars[cyl ==
cyl[i] ])),
                                                 replace = T),]
}

resampling_df <- do.call("rbind", resampling)
mtcars_oversample <- rbind(mtcars, resampling_df)
```

I get 307 observations. Anyone knows what I am doing wrong?

r sample resampling

asked Oct 15 '18 at 22:11

user3507584

2,260 ● 5 ● 24 ● 47

3 Answers

I think in this case, you can do the the sampling within groups using `data.table` 's `by=` functionality. `sample` from the `.I` row counter within each `cyl` group, and then use this

```
mtcars[mtcars[, sample(.I, 100, replace=TRUE), by=cyl]$V1,]
#      mpg cyl  disp  hp drat    wt  qsec vs am gear carb
#  1: 18.1   6 225.0 105 2.76 3.460 20.22  1  0    3    1
#  2: 17.8   6 167.6 123 3.92 3.440 18.90  1  0    4    4
#  3: 19.2   6 167.6 123 3.92 3.440 18.30  1  0    4    4
#  4: 19.2   6 167.6 123 3.92 3.440 18.30  1  0    4    4
#  5: 21.0   6 160.0 110 3.90 2.620 16.46  0  1    4    4
# ---
#296: 15.5   8 318.0 150 2.76 3.520 16.87  0  0    3    2
#297: 19.2   8 400.0 175 3.08 3.845 17.05  0  0    3    2
#298: 19.2   8 400.0 175 3.08 3.845 17.05  0  0    3    2
#299: 14.3   8 360.0 245 3.21 3.570 15.84  0  0    3    4
#300: 15.2   8 275.8 180 3.07 3.780 18.00  0  0    3    3
```

If you need to specify different counts for each group, here's one option. The special `.BY` object stores the value of the `by=` argument as a list.

```
grpcnt <- setNames(c(50,100,70), unique(mtcars$cyl))
#  6   4   8
# 50 100  70
mtcars[mtcars[, sample(.I, grpcnt[as.character(.BY[[1]])], replace=TRUE), by=cyl]$V1]
```

edited Oct 16 '18 at 22:47

answered Oct 15 '18 at 22:26

thelatemail

74.9k ● 10 ● 98 ● 162

I think this is the best and fastest solution. However, I didn't mention that I would need a different number of observations per cyl group. I guess I would use this solution stackoverflow.com/questions/33495916/… – user3507584 Oct 16 '18 at 6:38

1 @user3507584 - see my edit for how to adapt different group sizes to this sort of solution. – thelatemail Oct 16 '18 at 22:47

For an alternative solution, you can use `dplyr` and do:

4

```
library(dplyr)

mtcars %>%
    group_by(cyl) %>%
    do(sampled = sample_n(., size = 100, replace = TRUE)) %>%
    select(-cyl) %>%
    unnest()
```

answered Oct 15 '18 at 22:20

Marius

43.3k ● 10 ● 84 ● 86

3

Here's another way using `dplyr::slice`

```
mtcars %>%
  group_by(cyl) %>%
  slice(sample(n(), 100, replace = T)) %>%
  ungroup()
```

answered Oct 15 '18 at 22:25

Shree

9,243 ● 1 ● 8 ● 28
