

Logistic Regression example

This post uses the same German credit data that is used in the book.

A logistic regression model is fitted to the data to predict default.

##	Duration	Amount	InstallmentRatePercentage	ResidenceDuration	Age
## 1	6	1169	4	4	67
## 2	48	5951	2	2	22
## 3	12	2096	2	3	49
## 4	42	7882	2	4	45
## 5	24	4870	3	4	53
## 6	36	9055	2	4	35
##	NumberExistingCredits	NumberPeopleMaintenance	Telephone	Foreign	Worker
## 1	2	1	0		

```

1
## 2          1          1          1
1
## 3          1          2          1
1
## 4          1          2          1
1
## 5          2          2          1
1
## 6          1          2          0
1
##      Class  CheckingAccountStatus.lt.0  CheckingAccountStatus.0.to.200
## 1  Good          1          0
## 2  Bad           0          1
## 3  Good          0          0
## 4  Good          1          0
## 5  Bad           1          0
## 6  Good          0          0
##      CheckingAccountStatus.gt.200  CheckingAccountStatus.none
## 1          0          0
## 2          0          0
## 3          0          1
## 4          0          0
## 5          0          0
## 6          0          1
##      CreditHistory.NoCredit.AllPaid  CreditHistory.ThisBank.AllPaid
## 1          0          0
## 2          0          0
## 3          0          0
## 4          0          0
## 5          0          0
## 6          0          0
##      CreditHistory.PaidDuly  CreditHistory.Delay  CreditHistory.Critic
al
## 1          0          0
1
## 2          1          0
0
## 3          0          0
1
## 4          1          0
0

```

## 5	0	1		
0				
## 6	1	0		
0				
##	Purpose.NewCar	Purpose.UsedCar	Purpose.Furniture.Equipment	
## 1	0	0	0	
## 2	0	0	0	
## 3	0	0	0	
## 4	0	0	1	
## 5	1	0	0	
## 6	0	0	0	
##	Purpose.Radio.Television	Purpose.DomesticAppliance	Purpose.Repairs	
## 1	1	0		
0				
## 2	1	0		
0				
## 3	0	0		
0				
## 4	0	0		
0				
## 5	0	0		
0				
## 6	0	0		
0				
##	Purpose.Education	Purpose.Vacation	Purpose.Retaining	Purpose.Business
## 1	0	0	0	
0				
## 2	0	0	0	
0				
## 3	1	0	0	
0				
## 4	0	0	0	
0				
## 5	0	0	0	
0				
## 6	1	0	0	
0				
##	Purpose.Other	SavingsAccountBonds.lt.100	SavingsAccountBonds.100.to.500	
## 1	0	0		

0		
## 2	0	1
0		
## 3	0	1
0		
## 4	0	1
0		
## 5	0	1
0		
## 6	0	0
0		
## SavingsAccountBonds.500.to.1000	SavingsAccountBonds.gt.1000	
## 1	0	0
## 2	0	0
## 3	0	0
## 4	0	0
## 5	0	0
## 6	0	0
## SavingsAccountBonds.Unknown	EmploymentDuration.lt.1	
## 1	1	0
## 2	0	0
## 3	0	0
## 4	0	0
## 5	0	0
## 6	1	0
## EmploymentDuration.1.to.4	EmploymentDuration.4.to.7	
## 1	0	0
## 2	1	0
## 3	0	1
## 4	0	1
## 5	1	0
## 6	1	0
## EmploymentDuration.gt.7	EmploymentDuration.Unemployed	
## 1	1	0
## 2	0	0
## 3	0	0
## 4	0	0
## 5	0	0
## 6	0	0
## Personal.Male.Divorced.Seperated	Personal.Female.NotSingle	
## 1	0	0
## 2	0	1

##	3		0	0
##	4		0	0
##	5		0	0
##	6		0	0
##		Personal.Male.Single	Personal.Male.Married.Widowed	
##	1	1	0	
##	2	0	0	
##	3	1	0	
##	4	1	0	
##	5	1	0	
##	6	1	0	
##		Personal.Female.Single	OtherDebtorsGuarantors.None	
##	1	0	1	
##	2	0	1	
##	3	0	1	
##	4	0	0	
##	5	0	1	
##	6	0	1	
##		OtherDebtorsGuarantors.CoApplicant	OtherDebtorsGuarantors.Guara	
##	1		0	
	0			
##	2		0	
	0			
##	3		0	
	0			
##	4		0	
	1			
##	5		0	
	0			
##	6		0	
	0			
##		Property.RealEstate	Property.Insurance	Property.CarOther
##	1	1	0	0
##	2	1	0	0
##	3	1	0	0
##	4	0	1	0
##	5	0	0	0
##	6	0	0	0
##		Property.Unknown	OtherInstallmentPlans.Bank	OtherInstallmentPla
		ns.Stores		
##	1	0		0

0			
## 2	0	0	
0			
## 3	0	0	
0			
## 4	0	0	
0			
## 5	1	0	
0			
## 6	1	0	
0			
##	OtherInstallmentPlans.None	Housing.Rent	Housing.Own
Free			Housing.For
## 1	1	0	1
0			
## 2	1	0	1
0			
## 3	1	0	1
0			
## 4	1	0	0
1			
## 5	1	0	0
1			
## 6	1	0	0
1			
##	Job.UnemployedUnskilled	Job.UnskilledResident	Job.SkilledEmploy
ee			
## 1	0	0	
1			
## 2	0	0	
1			
## 3	0	1	
0			
## 4	0	0	
1			
## 5	0	0	
1			
## 6	0	1	
0			
##	Job.Management.SelfEmp.HighlyQualified		
## 1		0	
## 2		0	

```
## 3 0
## 4 0
## 5 0
## 6 0
```

```
with(GermanCredit, table(Class))
```

```
## Class
##   Bad Good
## 300  700
```

Split the data.

```
Train <- createDataPartition(GermanCredit$Class, p=0.6, list=FALSE)
training <- GermanCredit[ Train, ]
testing <- GermanCredit[ -Train, ]
```

Fit the logistic regression model, that is a GLM using a binomial link, using the caret function train().

```
mod_fit <- train(Class ~ Age + ForeignWorker + Property.RealEstate +
  Housing.Own +
  CreditHistory.Critical, data=training, method="g
  lm", family="binomial")
```

Transform the coefficients from log-odds to odds.

```
exp(coef(mod_fit$finalModel))
```

```
##           (Intercept)           Age           ForeignWor
ker
##           3.1475551           1.0119295           0.2162
119
##   Property.RealEstate   Housing.Own CreditHistory.Criti
cal
##           1.4981725           1.9448033           2.5952
555
```

```
predict(mod_fit, newdata=testing)
```

[illegible]


```

Good Good
## [267] Good Good Good Good Good Good Good Good Good Good Good Good Good
Good Good
## [281] Good Good Bad Good Good Good Good Good Good Good Good Good Good Good
Good Good
## [295] Good Good Good Good Good Good Good Good Good Good Good Good Good Good
Good Good
## [309] Bad Good Good Good Good Good Good Good Good Good Good Good Good Good
Good Good
## [323] Good Good Good Good Good Good Good Good Good Good Good Bad Good Good
Good Good
## [337] Good Good Good Good Good Bad Good Good Good Good Good Good Good Good
Good Good
## [351] Good Good Good Good Bad Good Good Good Good Good Good Good Good Good
Good Good
## [365] Good Good Good Good Good Good Good Good Good Good Good Good Good Good
Good Good
## [379] Bad Good Good Good Good Bad Good Good Good Good Good Good Good Good
Good Good
## [393] Good Good Good Good Good Good Bad Good
## Levels: Bad Good

```

```

predict(mod_fit, newdata=testing, type="prob")

```

```

##           Bad           Good
## 3    0.09802981 0.9019702
## 7    0.28724245 0.7127576
## 8    0.49245237 0.5075476
## 10   0.17278366 0.8272163
## 11   0.52208479 0.4779152
## 14   0.12504352 0.8749565
## 16   0.34078975 0.6592102
## 17   0.13441196 0.8655880
## 20   0.34345890 0.6565411
## 22   0.36791560 0.6320844
## 26   0.24759906 0.7524009
## 29   0.25204396 0.7479560
## 33   0.34613797 0.6538620
## 34   0.22360599 0.7763940
## 35   0.33813067 0.6618693
## 37   0.28161960 0.7183804

```

##	39	0.32759747	0.6724025
##	41	0.34613797	0.6538620
##	46	0.16123893	0.8387611
##	47	0.32239447	0.6776055
##	48	0.52799932	0.4720007
##	49	0.15492629	0.8450737
##	50	0.26569274	0.7343073
##	55	0.42773712	0.5722629
##	57	0.28967649	0.7103235
##	60	0.30120414	0.6987959
##	62	0.09698623	0.9030138
##	65	0.35695084	0.6430492
##	67	0.34882684	0.6511732
##	69	0.48652559	0.5134744
##	71	0.34613797	0.6538620
##	72	0.24706523	0.7529348
##	74	0.10674414	0.8932559
##	77	0.33548177	0.6645182
##	80	0.36791142	0.6320886
##	86	0.12109186	0.8789081
##	89	0.34613797	0.6538620
##	93	0.16447239	0.8355276
##	94	0.30874490	0.6912551
##	96	0.42483683	0.5751632
##	97	0.12375182	0.8762482
##	99	0.11253113	0.8874689
##	100	0.48948861	0.5105114
##	101	0.37622731	0.6237727
##	104	0.27212302	0.7278770
##	105	0.51912509	0.4808749
##	110	0.24981492	0.7501851
##	112	0.52799932	0.4720007
##	116	0.10014703	0.8998530
##	117	0.34613797	0.6538620
##	119	0.18142373	0.8185763
##	120	0.15964155	0.8403585
##	121	0.27269139	0.7273086
##	125	0.41591460	0.5840854
##	127	0.23886824	0.7611318
##	133	0.35423341	0.6457666
##	136	0.11018415	0.8898158
##	137	0.35423341	0.6457666

##	140	0.46582255	0.5341775
##	142	0.34613797	0.6538620
##	144	0.36791142	0.6320886
##	147	0.15492629	0.8450737
##	148	0.13719543	0.8628046
##	150	0.14436649	0.8556335
##	154	0.28644273	0.7135573
##	155	0.48948861	0.5105114
##	157	0.02322681	0.9767732
##	159	0.32499057	0.6750094
##	160	0.11492164	0.8850784
##	164	0.39048929	0.6095107
##	165	0.33021505	0.6697850
##	167	0.33813067	0.6618693
##	168	0.28460849	0.7153915
##	170	0.11859020	0.8814098
##	172	0.27034575	0.7296543
##	175	0.35695084	0.6430492
##	184	0.09698623	0.9030138
##	185	0.15964155	0.8403585
##	188	0.20177730	0.7982227
##	190	0.33813067	0.6618693
##	192	0.49541666	0.5045833
##	194	0.41016456	0.5898354
##	196	0.16284919	0.8371508
##	200	0.31980929	0.6801907
##	203	0.35695084	0.6430492
##	204	0.53390600	0.4660940
##	215	0.13719543	0.8628046
##	220	0.26129220	0.7387078
##	221	0.27034575	0.7296543
##	227	0.35423341	0.6457666
##	232	0.44523658	0.5547634
##	234	0.27269139	0.7273086
##	235	0.10561865	0.8943813
##	237	0.36515794	0.6348421
##	238	0.41616771	0.5838323
##	239	0.16123893	0.8387611
##	240	0.32239447	0.6776055
##	243	0.52504293	0.4749571
##	246	0.27269139	0.7273086
##	247	0.09492878	0.9050712

##	249	0.27034575	0.7296543
##	252	0.30453343	0.6954666
##	253	0.36241325	0.6375868
##	260	0.08995525	0.9100448
##	262	0.28241119	0.7175888
##	264	0.22360599	0.7763940
##	266	0.32759747	0.6724025
##	270	0.35967750	0.6403225
##	275	0.33548177	0.6645182
##	276	0.25654144	0.7434586
##	277	0.17620000	0.8238000
##	278	0.29705133	0.7029487
##	279	0.34078975	0.6592102
##	280	0.51023879	0.4897612
##	281	0.30120414	0.6987959
##	286	0.24981492	0.7501851
##	291	0.07416762	0.9258324
##	295	0.14436649	0.8556335
##	296	0.35695084	0.6430492
##	298	0.08742895	0.9125710
##	300	0.16610855	0.8338915
##	301	0.21000096	0.7899990
##	303	0.11135223	0.8886478
##	305	0.25150375	0.7484963
##	319	0.16610855	0.8338915
##	321	0.17278366	0.8272163
##	322	0.34078975	0.6592102
##	325	0.11253113	0.8874689
##	328	0.33548177	0.6645182
##	329	0.34345890	0.6565411
##	330	0.35152537	0.6484746
##	331	0.18873732	0.8112627
##	332	0.11983539	0.8801646
##	340	0.34345890	0.6565411
##	342	0.51912509	0.4808749
##	347	0.12887110	0.8711289
##	348	0.52799932	0.4720007
##	352	0.10380012	0.8961999
##	354	0.51320189	0.4867981
##	356	0.27742065	0.7225793
##	358	0.35695084	0.6430492
##	359	0.26569274	0.7343073

##	362	0.32759747	0.6724025
##	364	0.53390600	0.4660940
##	366	0.33021505	0.6697850
##	367	0.17448523	0.8255148
##	369	0.31467210	0.6853279
##	371	0.24759906	0.7524009
##	372	0.16447239	0.8355276
##	373	0.22342518	0.7765748
##	376	0.48652559	0.5134744
##	378	0.48948861	0.5105114
##	380	0.13579769	0.8642023
##	381	0.32239447	0.6776055
##	384	0.27034575	0.7296543
##	385	0.35695084	0.6430492
##	386	0.17792798	0.8220720
##	401	0.24103099	0.7589690
##	403	0.35423341	0.6457666
##	406	0.36791142	0.6320886
##	409	0.35695084	0.6430492
##	412	0.16123893	0.8387611
##	421	0.53685588	0.4631441
##	424	0.25204396	0.7479560
##	425	0.35967750	0.6403225
##	427	0.17109524	0.8289048
##	434	0.11613339	0.8838666
##	435	0.27269139	0.7273086
##	440	0.27034575	0.7296543
##	441	0.32239447	0.6776055
##	442	0.34613797	0.6538620
##	450	0.26821727	0.7317827
##	452	0.26338550	0.7366145
##	453	0.33548177	0.6645182
##	455	0.34345890	0.6565411
##	457	0.19781885	0.8021811
##	458	0.49245237	0.5075476
##	459	0.26801284	0.7319872
##	460	0.34078975	0.6592102
##	465	0.33548177	0.6645182
##	466	0.26358767	0.7364123
##	470	0.33284319	0.6671568
##	473	0.35152537	0.6484746
##	475	0.16447239	0.8355276

##	477	0.36241325	0.6375868
##	478	0.35967750	0.6403225
##	480	0.14732102	0.8526790
##	481	0.12887110	0.8711289
##	482	0.35695084	0.6430492
##	483	0.42773712	0.5722629
##	487	0.28967649	0.7103235
##	489	0.12282616	0.8771738
##	492	0.47172882	0.5282712
##	498	0.15338001	0.8466200
##	501	0.35423341	0.6457666
##	502	0.47172882	0.5282712
##	503	0.22000712	0.7799929
##	505	0.52504293	0.4749571
##	510	0.30705088	0.6929491
##	511	0.27034575	0.7296543
##	512	0.50134600	0.4986540
##	524	0.36515794	0.6348421
##	525	0.27034575	0.7296543
##	526	0.34613797	0.6538620
##	527	0.16775772	0.8322423
##	528	0.10561865	0.8943813
##	530	0.23671884	0.7632812
##	531	0.34078975	0.6592102
##	536	0.27684618	0.7231538
##	538	0.15805701	0.8419430
##	543	0.34345890	0.6565411
##	546	0.37067774	0.6293223
##	547	0.15805701	0.8419430
##	548	0.34078975	0.6592102
##	549	0.27504971	0.7249503
##	553	0.25204396	0.7479560
##	554	0.17448523	0.8255148
##	557	0.35152537	0.6484746
##	558	0.34882684	0.6511732
##	562	0.52504293	0.4749571
##	565	0.32759747	0.6724025
##	566	0.52799932	0.4720007
##	567	0.33021505	0.6697850
##	573	0.51912509	0.4808749
##	575	0.35423341	0.6457666
##	580	0.35423341	0.6457666

##	581	0.11983539	0.8801646
##	582	0.09802981	0.9019702
##	583	0.41879833	0.5812017
##	586	0.43621085	0.5637892
##	589	0.22410422	0.7758958
##	592	0.47172882	0.5282712
##	602	0.34613797	0.6538620
##	606	0.27980414	0.7201959
##	607	0.07475907	0.9252409
##	611	0.36791142	0.6320886
##	612	0.45404032	0.5459597
##	613	0.34882684	0.6511732
##	614	0.19662506	0.8033749
##	618	0.19594374	0.8040563
##	625	0.40468993	0.5953101
##	627	0.23671884	0.7632812
##	630	0.19100253	0.8089975
##	631	0.35152537	0.6484746
##	632	0.30957987	0.6904201
##	634	0.53980319	0.4601968
##	639	0.33548177	0.6645182
##	642	0.32499057	0.6750094
##	644	0.16447239	0.8355276
##	649	0.48356351	0.5164365
##	654	0.31467210	0.6853279
##	656	0.53095374	0.4690463
##	659	0.35152537	0.6484746
##	663	0.21797883	0.7820212
##	665	0.13860526	0.8613947
##	666	0.17448523	0.8255148
##	667	0.33548177	0.6645182
##	668	0.26801284	0.7319872
##	674	0.17966922	0.8203308
##	676	0.29376929	0.7062307
##	683	0.25428615	0.7457139
##	684	0.24486580	0.7551342
##	686	0.49541666	0.5045833
##	688	0.50727497	0.4927250
##	690	0.34345890	0.6565411
##	693	0.34882684	0.6511732
##	695	0.35152537	0.6484746
##	696	0.44816762	0.5518324

##	699	0.17966922	0.8203308
##	700	0.47764303	0.5223570
##	706	0.49245237	0.5075476
##	710	0.24539643	0.7546036
##	713	0.22617301	0.7738270
##	715	0.35423341	0.6457666
##	716	0.12120278	0.8787972
##	717	0.15338001	0.8466200
##	720	0.34345890	0.6565411
##	723	0.27504971	0.7249503
##	729	0.42194171	0.5780583
##	733	0.23035103	0.7696490
##	741	0.34078975	0.6592102
##	744	0.27980414	0.7201959
##	747	0.52799932	0.4720007
##	748	0.24539643	0.7546036
##	749	0.35695084	0.6430492
##	751	0.29705133	0.7029487
##	752	0.36515794	0.6348421
##	755	0.43355244	0.5664476
##	757	0.01716987	0.9828301
##	759	0.25880980	0.7411902
##	762	0.22137432	0.7786257
##	765	0.17448523	0.8255148
##	769	0.28887274	0.7111273
##	770	0.08429782	0.9157022
##	771	0.41879833	0.5812017
##	772	0.17792798	0.8220720
##	773	0.15964155	0.8403585
##	775	0.20562441	0.7943756
##	776	0.42168763	0.5783124
##	782	0.08708499	0.9129150
##	783	0.25880980	0.7411902
##	784	0.36515794	0.6348421
##	785	0.19843995	0.8015600
##	788	0.15648527	0.8435147
##	789	0.44816762	0.5518324
##	790	0.17448523	0.8255148
##	792	0.31723513	0.6827649
##	794	0.44523658	0.5547634
##	795	0.50134600	0.4986540
##	804	0.16123893	0.8387611

##	805	0.42458229	0.5754177
##	806	0.36241325	0.6375868
##	813	0.17966922	0.8203308
##	814	0.28482066	0.7151793
##	815	0.45992583	0.5400742
##	816	0.43646708	0.5635329
##	817	0.19469432	0.8053057
##	819	0.31212031	0.6878797
##	821	0.26801284	0.7319872
##	826	0.16941994	0.8305801
##	829	0.45698157	0.5430184
##	835	0.27269139	0.7273086
##	837	0.37067355	0.6293265
##	838	0.42748213	0.5725179
##	843	0.27742065	0.7225793
##	845	0.30202763	0.6979724
##	846	0.33284319	0.6671568
##	847	0.39614876	0.6038512
##	849	0.20033300	0.7996670
##	851	0.07644518	0.9235548
##	853	0.27212302	0.7278770
##	854	0.50134600	0.4986540
##	855	0.30705088	0.6929491
##	856	0.25428615	0.7457139
##	858	0.28887274	0.7111273
##	859	0.34882684	0.6511732
##	862	0.35152537	0.6484746
##	863	0.24981492	0.7501851
##	864	0.11735622	0.8826438
##	865	0.42168763	0.5783124
##	868	0.15032539	0.8496746
##	870	0.52504293	0.4749571
##	872	0.02377103	0.9762290
##	873	0.12492959	0.8750704
##	874	0.27504971	0.7249503
##	876	0.15338001	0.8466200
##	877	0.48948861	0.5105114
##	880	0.15964155	0.8403585
##	881	0.32499057	0.6750094
##	885	0.31212031	0.6878797
##	887	0.16284919	0.8371508
##	891	0.14881695	0.8511831

##	892	0.14436649	0.8556335
##	894	0.16284919	0.8371508
##	897	0.51320189	0.4867981
##	898	0.03990483	0.9600952
##	899	0.39873927	0.6012607
##	900	0.31467210	0.6853279
##	901	0.25374276	0.7462572
##	902	0.10339905	0.8966010
##	904	0.37901443	0.6209856
##	905	0.24759906	0.7524009
##	907	0.07581279	0.9241872
##	908	0.35423341	0.6457666
##	913	0.35967750	0.6403225
##	914	0.26569274	0.7343073
##	917	0.25654144	0.7434586
##	919	0.33813067	0.6618693
##	920	0.48060259	0.5193974
##	921	0.17278366	0.8272163
##	922	0.32759747	0.6724025
##	925	0.28241119	0.7175888
##	926	0.30453343	0.6954666
##	927	0.30621970	0.6937803
##	932	0.36791142	0.6320886
##	934	0.15032539	0.8496746
##	935	0.36515794	0.6348421
##	937	0.26569274	0.7343073
##	938	0.50727497	0.4927250
##	944	0.10561865	0.8943813
##	954	0.35695084	0.6430492
##	955	0.34882684	0.6511732
##	959	0.30705088	0.6929491
##	960	0.50727497	0.4927250
##	961	0.20935312	0.7906469
##	964	0.33284319	0.6671568
##	967	0.18142373	0.8185763
##	970	0.10788017	0.8921198
##	977	0.19100253	0.8089975
##	978	0.31467210	0.6853279
##	981	0.14002722	0.8599728
##	982	0.49838128	0.5016187
##	984	0.35695084	0.6430492
##	985	0.16941994	0.8305801

```
## 986 0.29623565 0.7037643
## 988 0.19100253 0.8089975
## 994 0.34613797 0.6538620
## 996 0.25880980 0.7411902
## 999 0.52799932 0.4720007
## 1000 0.17448523 0.8255148
```

Model Evaluation and Diagnostics

Fit two models with the R function `glm()`.

```
mod_fit_one <- glm(Class ~ Age + ForeignWorker + Property.RealEstate
+ Housing.Own +
                    CreditHistory.Critical, data=training, family="
binomial")
```

```
anova(mod_fit_one)
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Class
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev
## NULL              599      733.04
## Age                1      4.7017      598      728.34
## ForeignWorker      1      8.4069      597      719.93
## Property.RealEstate 1      5.0428      596      714.89
## Housing.Own        1     14.1010      595      700.78
## CreditHistory.Critical 1     17.6791      594      683.11
```

```
mod_fit_two <- glm(Class ~ Age + ForeignWorker, data=training, famil
y="binomial")
```

```
anova(mod_fit_two)
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Class
##
## Terms added sequentially (first to last)
##
##
##              Df Deviance Resid. Df Resid. Dev
## NULL                      599      733.04
## Age                1    4.7017      598      728.34
## ForeignWorker      1    8.4069      597      719.93
```

Goodness-of-Fit

Likelihood Ration Test

```
anova(mod_fit_one, mod_fit_two, test = "Chisq")
```

```
## Analysis of Deviance Table
##
## Model 1: Class ~ Age + ForeignWorker + Property.RealEstate + Hous
ing.Own +
##      CreditHistory.Critical
## Model 2: Class ~ Age + ForeignWorker
##      Resid. Df Resid. Dev Df Deviance  Pr(>Chi)
## 1          594      683.11
## 2          597      719.93 -3   -36.823 5.016e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
library(lmtest)
```

```
## Loading required package: zoo
```

```
##  
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':  
##  
##      as.Date, as.Date.numeric
```

```
lrtest(mod_fit_one, mod_fit_two)
```

```
## Likelihood ratio test  
##  
## Model 1: Class ~ Age + ForeignWorker + Property.RealEstate + Hous  
ing.Own +  
##      CreditHistory.Critical  
## Model 2: Class ~ Age + ForeignWorker  
##      #Df  LogLik Df  Chisq Pr(>Chisq)  
## 1      6 -341.55  
## 2      3 -359.96 -3  36.823  5.016e-08 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Pseudo R^2 .

```
library(pscl)
```

```
## Loading required package: MASS
```

```
## Classes and Methods for R developed in the
```

```
## Political Science Computational Laboratory
```

```
## Department of Political Science
```

```
## Stanford University
```

```
## Simon Jackman
```

```
## hurdle and zeroinfl functions by Achim Zeileis
```

```
pR2(mod_fit_one) # look for 'McFadden'
```

```
##              llh              llhNull              G2              McFadden
r2ML
## -341.55286427 -366.51858123    49.93143393    0.06811583    0.0798
5044
##              r2CU
##      0.11321811
```

Hosmer-Lemeshow Test

```
library(MKmisc)
HLgof.test(fit = fitted(mod_fit_one), obs = training$Class)
```

```
## Warning in Ops.factor(1, obs): '-' not meaningful for factors
```

```
## Warning in Ops.factor(1, obs): '-' not meaningful for factors
```

```
## $C
##
## Hosmer-Lemeshow C statistic
##
## data: fitted(mod_fit_one) and training$Class
## X-squared = 600, df = 8, p-value < 2.2e-16
##
##
## $H
##
## Hosmer-Lemeshow H statistic
##
## data: fitted(mod_fit_one) and training$Class
## X-squared = 600, df = 8, p-value < 2.2e-16
```

```
library(ResourceSelection)
```

```
## ResourceSelection 0.3-2    2017-02-28
```

```
hoslem.test(training$Class, fitted(mod_fit_one), g=10)
```

```
## Warning in Ops.factor(1, y): '-' not meaningful for factors
```

```
##
## Hosmer and Lemeshow goodness of fit (GOF) test
##
## data: training$Class, fitted(mod_fit_one)
## X-squared = 600, df = 8, p-value < 2.2e-16
```

Statistical Tests for Individual Predictors

Wald Test

```
library(survey)
```

```
## Loading required package: grid
```

```
## Loading required package: Matrix
```

```
## Loading required package: survival
```

```
##  
## Attaching package: 'survival'
```

```
## The following object is masked from 'package:caret':  
##  
##      cluster
```

```
##  
## Attaching package: 'survey'
```

```
## The following object is masked from 'package:graphics':  
##  
##      dotchart
```

```
regTermTest(mod_fit_one, "ForeignWorker")
```

```
## Wald test for ForeignWorker  
## in glm(formula = Class ~ Age + ForeignWorker + Property.RealEstate +  
##      Housing.Own + CreditHistory.Critical, family = "binomial",  
##      data = training)  
## F = 4.157408 on 1 and 594 df: p= 0.041895
```

Variable Importance


```
varImp(mod_fit)
```

```
## glm variable importance
##
## Overall
## CreditHistory.Critical 100.00
## Housing.Own             77.77
## ForeignWorker           24.55
## Property.RealEstate     16.12
## Age                     0.00
```

Validation of Predicted Values

Classification Rate

Accuracy.

```
pred = predict(mod_fit, newdata=testing)

accuracy <- table(pred, testing[, "Class"])
accuracy
```

```
##
## pred    Bad Good
##    Bad      9   19
##    Good 111  261
```

```
sum(diag(accuracy))/sum(accuracy)
```

```
## [1] 0.675
```

Confussion Matrix.

```
pred = predict(mod_fit, newdata=testing)
confusionMatrix(data=pred, testing$Class)
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Bad Good
##           Bad    9   19
##           Good 111  261
##
##           Accuracy : 0.675
##           95% CI : (0.6267, 0.7207)
##           No Information Rate : 0.7
##           P-Value [Acc > NIR] : 0.8736
##
##           Kappa : 0.0091
##           Mcnemar's Test P-Value : 1.449e-15
##
##           Sensitivity : 0.0750
##           Specificity : 0.9321
##           Pos Pred Value : 0.3214
##           Neg Pred Value : 0.7016
##           Prevalence : 0.3000
##           Detection Rate : 0.0225
##           Detection Prevalence : 0.0700
##           Balanced Accuracy : 0.5036
##
##           'Positive' Class : Bad
##

```

ROC

From the blog post, “The receiving operating characteristic is a measure of classifier performance. Using the proportion of positive data points that are correctly considered as positive and the proportion of negative data points that are mistakenly considered as positive, we generate a graphic that shows the trade off between the rate at which you can correctly predict something with the rate of incorrectly predicting something. Ultimately, we’re concerned about the area under the ROC curve, or AUROC. That metric ranges from 0.50 to 1.00, and values above 0.80 indicate that the model does a good job in discriminating between the two categories which comprise our target variable.”


```
library(ROCR)
```

```
## Loading required package: gplots
```

```
##  
## Attaching package: 'gplots'
```

```
## The following object is masked from 'package:stats':  
##  
##      lowess
```

```
# Compute AUC for predicting Class with the model  
prob <- predict(mod_fit_one, newdata=testing, type="response")  
pred <- prediction(prob, testing$Class)  
perf <- performance(pred, measure = "tpr", x.measure = "fpr")  
plot(perf)
```



```
auc <- performance(pred, measure = "auc")  
auc <- auc@y.values[[1]]  
auc
```

```
## [1] 0.6469196
```

K-Fold Cross Validation

Split the data into k folds.

```

ctrl <- trainControl(method = "repeatedcv", number = 10, savePredictions = TRUE)

mod_fit <- train(Class ~ Age + ForeignWorker + Property.RealEstate + Housing.Own +
                  CreditHistory.Critical, data=GermanCredit, method="glm", family="binomial",
                  trControl = ctrl, tuneLength = 5)

pred <- predict(mod_fit, newdata=testing)
confusionMatrix(data=pred, testing$Class)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction Bad Good
##           Bad    7   12
##           Good 113  268
##
##           Accuracy : 0.6875
##           95% CI : (0.6396, 0.7326)
##           No Information Rate : 0.7
##           P-Value [Acc > NIR] : 0.7273
##
##           Kappa : 0.0204
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.05833
##           Specificity : 0.95714
##           Pos Pred Value : 0.36842
##           Neg Pred Value : 0.70341
##           Prevalence : 0.30000
##           Detection Rate : 0.01750
##           Detection Prevalence : 0.04750
##           Balanced Accuracy : 0.50774
##
##           'Positive' Class : Bad
##

```

Randomly sample with replacement. B = 1000

```
start.time <- Sys.time()

ctrl <- trainControl(method = "boot632", number = 1000, savePredictions = TRUE)

mod_fit <- train(Class ~ Age + ForeignWorker + Property.RealEstate + Housing.Own +
                  CreditHistory.Critical, data=GermanCredit, method="glm", family="binomial",
                  trControl = ctrl, tuneLength = 5)

pred <- predict(mod_fit, newdata=testing)
confusionMatrix(data=pred, testing$Class)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction Bad Good
##           Bad    7   12
##           Good 113  268
##
##           Accuracy : 0.6875
##           95% CI : (0.6396, 0.7326)
##           No Information Rate : 0.7
##           P-Value [Acc > NIR] : 0.7273
##
##           Kappa : 0.0204
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.05833
##           Specificity : 0.95714
##           Pos Pred Value : 0.36842
##           Neg Pred Value : 0.70341
##           Prevalence : 0.30000
##           Detection Rate : 0.01750
##           Detection Prevalence : 0.04750
##           Balanced Accuracy : 0.50774
##
##           'Positive' Class : Bad
##
```

```
end.time <- Sys.time()
time.taken <- end.time - start.time
time.taken
```

```
## Time difference of 14.84216 secs
```

Using parallel processing.

```
start.time <- Sys.time()

library(doMC)
```

```
## Loading required package: foreach
```

```
## Loading required package: iterators
```

```
## Loading required package: parallel
```

```
registerDoMC(cores = 4)
```

```
ctrl <- trainControl(method = "boot632", number = 1000, savePredictions = TRUE,  
                     allowParallel = TRUE)
```

```
mod_fit <- train(Class ~ Age + ForeignWorker + Property.RealEstate +  
Housing.Own +  
                CreditHistory.Critical, data=GermanCredit, method="glm", family="binomial",  
                trControl = ctrl, tuneLength = 5)
```

```
pred <- predict(mod_fit, newdata=testing)  
confusionMatrix(data=pred, testing$Class)
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction Bad Good
##           Bad    7   12
##           Good 113  268
##
##           Accuracy : 0.6875
##           95% CI : (0.6396, 0.7326)
##           No Information Rate : 0.7
##           P-Value [Acc > NIR] : 0.7273
##
##           Kappa : 0.0204
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.05833
##           Specificity : 0.95714
##           Pos Pred Value : 0.36842
##           Neg Pred Value : 0.70341
##           Prevalence : 0.30000
##           Detection Rate : 0.01750
##           Detection Prevalence : 0.04750
##           Balanced Accuracy : 0.50774
##
##           'Positive' Class : Bad
##
```

```
end.time <- Sys.time()
time.taken <- end.time - start.time
time.taken
```

```
## Time difference of 6.356072 secs
```