



R news and tutorials contributed by hundreds of R bloggers

- [Home](#)
- [About](#)
- [RSS](#)
- [add your blog!](#)
- [Learn R](#)
- [R jobs](#) 
- [Contact us](#)

Welcome!

 Follow @rbloggers

80.2k followers

Here you will find daily **news and tutorials about R**, contributed by hundreds of bloggers. There are many ways to **follow us** -

[By e-mail:](#)


51451 readers

BY FEEDBURNER

[On Facebook:](#)



R blog...
77K likes

 Like Page

Be the first of your friends to like this

If you are an R blogger yourself you are invited to [add your own R content feed to this site](#) (Non-English R bloggers should add themselves- [here](#))

[Jobs for R-users](#)

- [\(Junior\) Data Analyst – Berlin](#)
- [Associate Researcher \(Data Analytics\)](#)
- [Lead Data Scientist @ Stabio, Ticino, Switzerland](#)

- [Senior Scientist, Translational Informatics](#)
- [Senior Scientist, Translational Informatics @ Vancouver, BC, Canada](#)

Recent Posts

- [Easy introduction to Offensive Programming](#)
- [Le Monde puzzle \[#1114\]](#)
- [upcoming AI-related courses](#)
- [Free R/datascience Extract: Evaluating a Classification Model with a Spam Filter](#)
- [Parsing Sda Pages](#)
- [Super Solutions for Shiny Apps #4 of 5: Using R6 Classes](#)
- [Simulating data with Bayesian networks](#)
- [JAMA retraction after miscoding – new Finalfit function to check recoding](#)
- [Finding free science books from Springer](#)
- [Merge MLP And CNN in Keras](#)
- [Shiny 1.4.0](#)
- [Strange Attractors: an R experiment about maths, recursivity and creative coding](#)
- [What are Your Use Cases for rOpenSci Tools and Resources?](#)
- [Selection bias, death, and dying](#)
- [Automatic data types checking in predictive models](#)

Other sites

- [Jobs for R-users](#)
- [SAS blogs](#)

How to perform a Logistic Regression in R

September 13, 2015
By [Michy Alice](#)

 Like 166

 Share

 Tweet

 Share

[This article was first published on [DataScience+](#), and kindly contributed to [R-bloggers](#).
(You can report issue about the content on this page [here](#))

Want to share your content on R-bloggers? [click here](#) if you have a blog, or [here](#) if you don't.

 Share

 Tweet

Logistic regression is a method for fitting a regression curve, $y = f(x)$, when y is a categorical variable. The typical use of this model is predicting y given a set of predictors x . The predictors can be continuous, categorical or a mix of both.

The categorical variable y , in general, can assume different values. In the simplest case scenario y is binary meaning that it can assume either the value 1 or 0. A classical example used in machine learning is email classification: given a set of attributes for each email such as number of words, links and pictures, the algorithm should decide whether the email is spam (1) or not (0). In this post we call the model “**binomial logistic regression**”, since the variable to predict is binary, however, logistic regression can also be used to predict a dependent variable which can assume more than 2 values. In this second case we call the model “multinomial logistic regression”. A typical example for instance, would be classifying films between “Entertaining”, “borderline” or “boring”.

Logistic regression implementation in R

R makes it very easy to fit a logistic regression model. The function to be called is `glm()` and the fitting process is not so different from the one used in linear regression. In this post I am going to fit a binary logistic regression model and explain each step.

The dataset

We’ll be working on the *Titanic dataset*. There are different versions of this datasets freely available online, however I suggest to use the one available at [Kaggle](#), since it is almost ready to be used (in order to download it you need to sign up to Kaggle). The dataset (training) is a collection of data about some of the passengers (889 to be precise), and the goal of the competition is to predict the survival (either 1 if the passenger survived or 0 if they did not) based on some features such as the *class of service*, the *sex*, the *age* etc. As you can see, we are going to use both categorical and continuous variables.

The data cleaning process

When working with a real dataset we need to take into account the fact that some data might be missing or corrupted, therefore we need to prepare the dataset for our analysis. As a first step we [load the csv data](#) using the `read.csv()` function.
Make sure that the parameter `na.strings` is equal to `c("")` so that each missing value is coded as a NA. This will help us in the next steps.

```
training.data.raw <- read.csv('train.csv',header=T,na.strings=c(""))
```

Now we need to check for missing values and look how many unique values there are for each variable using the `sapply()` function which applies the function passed as argument to each column of the dataframe.

```
sapply(training.data.raw,function(x) sum(is.na(x)))
```


PassengerId	Survived	Pclass	Name	Sex
0	0	0	0	0
Age	SibSp	Parch	Ticket	Fare
177	0	0	0	0
Cabin	Embarked			
687	2			

```
sapply(training.data.raw, function(x) length(unique(x)))
```

PassengerId	Survived	Pclass	Name	Sex
891	2	3	891	2
Age	SibSp	Parch	Ticket	Fare
89	7	7	681	248
Cabin	Embarked			
148	4			

A visual take on the missing values might be helpful: the `Amelia` package has a special plotting function `missmap()` that will plot your dataset and highlight missing values:

```
library(Amelia)
missmap(training.data.raw, main = "Missing values vs observed")
```



The variable `cabin` has too many missing values, we will not use it. We will also drop `PassengerId` since it is only an index and `Ticket`. Using the `subset()` function we [subset](#) the original dataset selecting the relevant columns only.

```
data <- subset(training.data.raw,select=c(2,3,5,6,7,8,10,12))
```

Taking care of the missing values

Now we need to account for the other missing values. R can easily deal with them when fitting a generalized linear model by setting a parameter inside the fitting function. However, personally I prefer to replace the NAs “by hand”, when is possible. There are different ways to do this, a typical approach is to replace the missing values with the average, the median or the mode of the existing one. I’ll be using the average.

```
data$Age[is.na(data$Age)] <- mean(data$Age,na.rm=T)
```

As far as categorical variables are concerned, using the `read.table()` or `read.csv()` by default will encode the categorical variables as factors. A factor is how R deals categorical variables. We can check the encoding using the following lines of code

```
is.factor(data$Sex)
TRUE

is.factor(data$Embarked)
TRUE
```

For a better understanding of how R is going to deal with the categorical variables, we can use the `contrasts()` function. This function will show us how the variables have been dummyfied by R and how to interpret them in a model.

```
contrasts(data$Sex)
      male
female    0
male      1

contrasts(data$Embarked)
      Q S
C 0 0
Q 1 0
S 0 1
```

For instance, you can see that in the variable sex, female will be used as the reference. As for the missing values in Embarked, since there are only two, we will discard those two rows (we could also have replaced the missing values with the mode and keep the datapoints).

```
data <- data[!is.na(data$Embarked),]
rownames(data) <- NULL
```

Before proceeding to the fitting process, let me remind you how important is *cleaning and formatting of the data*. This preprocessing step often is crucial for obtaining a good fit of the model and better predictive ability.

Model fitting

We split the data into two chunks: training and testing set. The training set will be used to fit our model which we will be testing over the testing set.

```
train <- data[1:800,]
test  <- data[801:889,]
```

Now, let’s fit the model. Be sure to specify the parameter `family=binomial` in the `glm()` function.

```
model <- glm(Survived ~.,family=binomial(link='logit'),data=train)
```

By using function `summary()` we obtain the results of our model:

```
summary(model)

Call:
glm(formula = Survived ~ ., family = binomial(link = "logit"),
    data = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max 
-2.6064  -0.5954  -0.4254   0.6220   2.4165 

Coefficients:
      Estimate Std. Error z value Pr(>|z|)
(Intercept)  5.137627   0.594998   8.635  < 2e-16 ***
Pclass      -1.087156   0.151168  -7.192 6.40e-13 ***
Sexmale     -2.756819   0.212026 -13.002  < 2e-16 ***
Age         -0.037267   0.008195  -4.547 5.43e-06 ***
SibSp       -0.292920   0.114642  -2.555  0.0106 *
Parch       -0.116576   0.128127  -0.910  0.3629
```

```
Fare      0.001528    0.002353    0.649    0.5160
EmbarkedQ -0.002656    0.400882   -0.007    0.9947
EmbarkedS -0.318786    0.252960   -1.260    0.2076
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 1065.39 on 799 degrees of freedom
Residual deviance: 709.39 on 791 degrees of freedom
AIC: 727.39
```

Number of Fisher Scoring iterations: 5

Interpreting the results of our logistic regression model

Now we can analyze the fitting and interpret what the model is telling us.

First of all, we can see that *SibSp*, *Fare* and *Embarked* are not statistically significant. As for the statistically significant variables, sex has the lowest p-value suggesting a strong association of the sex of the passenger with the probability of having survived. The negative coefficient for this predictor suggests that all other variables being equal, the male passenger is less likely to have survived. Remember that in the logit model the response variable is log odds: $\ln(\text{odds}) = \ln(p/(1-p)) = a \cdot x_1 + b \cdot x_2 + \dots + z \cdot x_n$. Since male is a dummy variable, being male reduces the log odds by 2.75 while a unit increase in age reduces the log odds by 0.037.

Now we can run the `anova()` function on the model to analyze the table of deviance

```
anova(model, test="Chisq")
```

```
Analysis of Deviance Table
Model: binomial, link: logit
Response: Survived
Terms added sequentially (first to last)
```

	Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
NULL			799	1065.39	
Pclass	1	83.607	798	981.79	< 2.2e-16 ***
Sex	1	240.014	797	741.77	< 2.2e-16 ***
Age	1	17.495	796	724.28	2.881e-05 ***
SibSp	1	10.842	795	713.43	0.000992 ***
Parch	1	0.863	794	712.57	0.352873
Fare	1	0.994	793	711.58	0.318717
Embarked	2	2.187	791	709.39	0.334990

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The difference between the null deviance and the residual deviance shows how our model is doing against the null model (a model with only the intercept). The wider this gap, the better. Analyzing the table we can see the drop in deviance when adding each variable one at a time. Again, adding *Pclass*, *Sex* and *Age* significantly reduces the residual deviance. The other variables seem to improve the model less even though *SibSp* has a low p-value. A large p-value here indicates that the model without the variable explains more or less the same amount of variation. Ultimately what you would like to see is a significant drop in deviance and the AIC.

While no exact equivalent to the R^2 of linear regression exists, the McFadden R^2 index can be used to assess the model fit.

```
library(pscl)
pR2(model)

      llh      llhNull      G2      McFadden      r2ML      r2CU
-354.6950111 -532.6961008 356.0021794  0.3341513  0.3591775  0.4880244
```

Assessing the predictive ability of the model

In the steps above, we briefly evaluated the fitting of the model, now we would like to see how the model is doing when predicting y on a new set of data. By setting the parameter `type='response'`, R will output probabilities in the form of $P(y=1|X)$. Our decision boundary will be 0.5. If $P(y=1|X) > 0.5$ then $y = 1$ otherwise $y=0$. Note that for some applications different thresholds could be a better option.

```
fitted.results <- predict(model,newdata=subset(test,select=c(2,3,4,5,6,7,8)),type='response')
fitted.results <- ifelse(fitted.results > 0.5,1,0)
```

```
misClasificError <- mean(fitted.results != test$Survived)
print(paste('Accuracy',1-misClasificError))
```

```
"Accuracy 0.842696629213483"
```

The 0.84 accuracy on the test set is quite a good result. However, keep in mind that this result is somewhat dependent on the manual split of the data that I made earlier, therefore if you wish for a more precise score, you would be better off running some kind of cross validation such as k-fold cross validation.

As a last step, we are going to plot the *ROC curve* and calculate the *AUC* (area under the curve) which are typical performance measurements for a binary classifier.


The ROC is a curve generated by plotting the true positive rate (TPR) against the false positive rate (FPR) at various threshold settings while the AUC is the area under the ROC curve. As a rule of thumb, a model with good predictive ability should have an AUC closer to 1 (1 is ideal) than to 0.5.

```
library(ROCR)
p <- predict(model, newdata=subset(test,select=c(2,3,4,5,6,7,8)), type="response")
pr <- prediction(p, test$Survived)
prf <- performance(pr, measure = "tpr", x.measure = "fpr")
plot(prf)
```

```
auc <- performance(pr, measure = "auc")
auc <- auc@y.values[[1]]
auc
```

```
0.8647186
```

And here is the ROC plot:



I hope this post will be useful. A gist with the full code for this example can be found [here](#).

Thank you for reading this post, leave a comment below if you have any question.

◆ 10 comments on this item ◆ Share on Facebook ◆ Stumble It! ◆ Digg This!

 Share

 Tweet

To **leave a comment** for the author, please follow the link and comment on their blog:
[DataScience+.](#)

R-bloggers.com offers **daily e-mail updates** about [R](#) news and tutorials about [learning R](#) and many other topics. [Click here if you're looking to post or find an R/data-science job.](#)

Want to share your content on R-bloggers? [click here](#) if you have a blog, or [here](#) if you don't.

If you got this far, why not **subscribe for updates** from the site?
Choose your flavor: [e-mail](#), [twitter](#), [RSS](#), or [facebook](#)...

 Like 166

 Share

 Tweet

 Share

Comments are closed.

Search R-bloggers

Go

Most visited articles of the week

1. [How to write the first for loop in R](#)
2. [5 Ways to Subset a Data Frame in R](#)
3. [R – Sorting a data frame by the contents of a column](#)
4. [Loading packages efficiently](#)
5. [Using apply, sapply, lapply in R](#)
6. [In-depth introduction to machine learning in 15 hours of expert videos](#)
7. [R Tutorial Series: Simple Linear Regression](#)
8. [Date Formats in R](#)
9. [Installing R packages](#)

Sponsors

Submit your abstract for the Enterprise Applications of the R Language Conference (EARL) before 31/3/20

London 8-10 September, 2020




Learn **R**
by doing.



plotly | Dash for R

Deliver Discovery
with interactive analytic apps

Download the white paper



MACHINE
LEARNING
WEEK


Workshop
Machine Learning
with R: A Hands-On
Introduction

Robert Muenchen
University of Tennessee



TRY **R** Studio Team

DOWNLOAD QUICKSTART VM



Beginner's Guide to
Spatial, Temporal and
Spatial-Temporal Ecological
Data Analysis with R-INLA

Zuur, Ieno, Saveliev



QUANTIDE
knowledge from data

Quantide: statistical consulting and training

Apply **[R]**
Programming

Build Applications with Shiny

40% Off Deal
Ends Friday



ODSC
BOSTON

TRAININGS: R & PYTHON
Einführung & Machine Learning
BERLIN HAMBURG
April 2020 Oktober 2020
[Jetzt anmelden](#)



jumping rivers

TRAINING: R, SCALA, STAN

STATWORX

Data Science Service

Data Science | Consulting | Development | Training

Try the FASTEST ML
for **R**



Click for a Free Trial



ALPHIEN

BUILD QUANT STRATEGIES
JOIN COMPETITIONS
EARN LICENSING FEES

[VISIT THE QUANT LAB](#)

[Contact us](#) if you wish to help support R-bloggers, and place **your banner here**.

Submit your abstract for the Enterprise Applications of the R Language Conference (EARL) before 31/3/20

London 8-10 September, 2020


EARLconference

MANGOSOLUTIONS

DataCamp

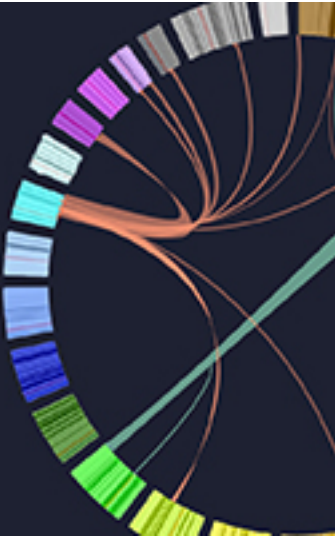
Learn **R** by doing.




plotly | Dash for R


Deliver Discovery with interactive analytic apps

Download the white paper



MACHINE LEARNING WEEK

Workshop Machine Learning with R: A Hands-On Introduction



Robert Muenchen
University of Tennessee

TRY **R** Studio Team

DOWNLOAD QUICKSTART VM



Beginner's Guide to

Spatial, Temporal and Spatial-Temporal Ecological Data Analysis with R-INLA

Zuur, Ieno, Saveliev



Quantide: statistical consulting and training

Apply [R] Programming

Build Applications with Shiny

40% Off Deal
Ends Friday



TRAININGS: R & PYTHON

Einführung & Machine Learning

BERLIN **HAMBURG**

April 2020 Oktober 2020


[Jetzt anmelden](#)



Become a "Barefoot Writer"

And live the life of a famous writer without having to write a block-buster movie script or a chart-topping book.

LEARN MORE ABOUT THIS UNIQUE GROUP OF HIGHLY-PAID WRITERS ...



jumping rivers

TRAINING: R, SCALA, STAN

STATWORX

Data Science Service

Data Science | Consulting | Development | Training

Try the FASTEST ML for R



Click for a Free Trial





ALPHIEN

BUILD QUANT STRATEGIES
JOIN COMPETITIONS
EARN LICENSING FEES

VISIT THE QUANT LAB

Our ads respect your privacy. Read our [Privacy Policy page](#) to learn more.

[Contact us](#) if you wish to help support R-bloggers, and place **your banner here**.

[Jobs for R users](#)

- [\(Junior\) Data Analyst – Berlin](#)
- [Associate Researcher \(Data Analytics\)](#)
- [Lead Data Scientist @ Stabio, Ticino, Switzerland](#)
- [Senior Scientist, Translational Informatics](#)
- [Senior Scientist, Translational Informatics @ Vancouver, BC, Canada](#)
- [Data Scientist @ New York, United States](#)
- [Empirical Research Librarian](#)

[Full list of contributing R-bloggers](#)

[R-bloggers](#) was founded by [Tal Galili](#), with gratitude to the [R](#) community.

Is powered by [WordPress](#) using a [bavotasan.com](#) design.

Copyright © 2019 **R-bloggers**. All Rights Reserved. [Terms and Conditions](#) for this website

